



Investment Association Presentation

Machine Learning in Algorithm Trading:
An Example of Recurrent Reinforcement Learning
(Partial Achievement of CS534 Artificial Intelligence Project)

Ethan Song
April 11, 2016

Concepts



- Machine Learning (ML):
 - Develop the computer's ability of learning without explicitly programmed
 - E.g. Facebook face recognition
- Algorithm Trading (AT):
 - The trading strategies are programmed and implemented by computers
 - Previous work: Emerging from 1980s, predictive math models, statistical arbitrage, event-driven system have been using for decades
 - New method: application of machine learning

Concepts



- Reinforcement Learning (RL):
 - Example: Chess game learning
 - Model Components:
 - I. Environment states $\{S_i\}_n$
 - II. Actions $\{A_i\}_n$
 - III. Transition Model: How does the state change from step i to step $i+1$?
e.g. weather change model, transitional probability matrix
 - IV. Reward $\{R_i\}_n$: immediate step reward
 - V. Utilities: A goal function regarding total rewards

RRL: Initialization



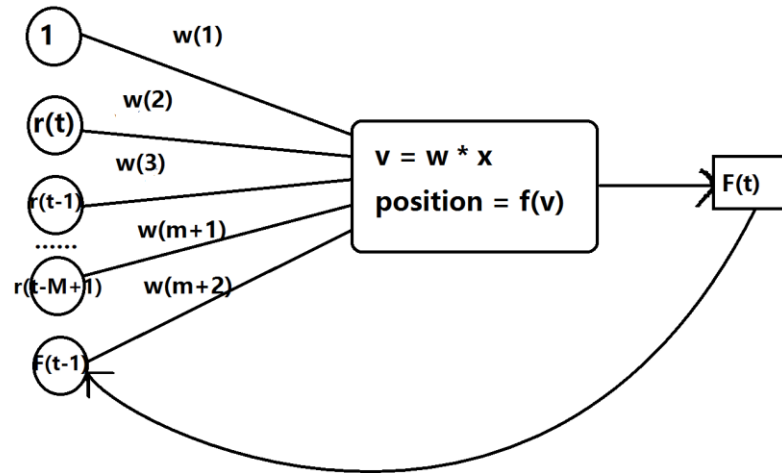
- Sample Data:
 - Simplified Data: Time and close price
 - Assumptions & Denotations:
 - Only trade one share per transaction
 - P_t : price at date t
 - r_t : price increment at date t , $r_t = P_t - P_{t-1}$
 - N : dimension of observations (number of the observed previous dates)
 - w : weights, $w \in \mathbb{R}^{N+2}$ (explained later)
 - u : commission fees per trade
 - T : length of time period
- Model Components:
 - States x_t , defined as the vector:
 - $[1, r_t, r_{t-1}, r_{t-2}, \dots, r_{t-N+1}, F_{t-1}]$
 - e.g. For 2014/9/5, if $N = 4$, $x_t = [1, 5, 5, -10, -5, 1]$ as showed in the sample data with a previous long position F_{t-1} denoted as $F_{t-1} = 1$ on 2014/9/4.
 - Actions/Positions F_t : The position on date t . F_t can be 1 as a long position, 0 as no position, -1 as a short position.
 - Immediate Rewards: $R_t = F_{t-1} * r_t - u * \text{abs}(F_t - F_{t-1})$
 - Utility: $S_T = E[R_t] / \text{Std}(R_t) = E[R_t] / \sqrt{E[R_t^2] - (E[R_t])^2}$ “profit per unit risk”
 - Denote $A = E[R_t]$, $B = E[R_t^2]$, then $S_T = A / \sqrt{B - A^2}$ (Similar to Sharp Ratio)

2014/9/1	10740
2014/9/2	10735
2014/9/3	10725
2014/9/4	10730
2014/9/5	10735


RRL: Model Description



- Neuro Structure:



- Input: $v = w^T * x = w_1 + w_2 r_t + w_3 r_{t-1} + \dots + w_{N+1} r_{t-M+1} + w_{N+2} F_{t-1}$
- $f(v) = \tanh(v) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$ (Hyperbolic Tangent)



RRL: Mathematical Principle



- $w^{i+1} = w^i + \alpha * (dS_T / dw^i)$ (online learning, gradient ascent)
- How to calculate (dS_T / dw^i) ?

Recall $\frac{dS_T}{dw} = \frac{d}{dw} \left\{ \frac{A}{\sqrt{B-A^2}} \right\}$ where $A = E(R_t) = \frac{1}{T} \sum_{t=1}^T R_t$ $B = E(R_t^2) = \frac{1}{T} \sum_{t=1}^T R_t^2$

$$\begin{aligned} \frac{dS_T}{dw} &= \frac{d}{dw} \left\{ \frac{A}{\sqrt{B-A^2}} \right\} = \frac{dS_T}{dA} \frac{dA}{dw} + \frac{dS_T}{dB} \frac{dB}{dw} = \sum_{t=1}^T \left(\frac{dS_T}{dA} \frac{dA}{dR_t} + \frac{dS_T}{dB} \frac{dB}{dR_t} \right) \frac{dR_t}{dw} \\ &= \sum_{t=1}^T \left(\frac{dS_T}{dA} \frac{dA}{dR_t} + \frac{dS_T}{dB} \frac{dB}{dR_t} \right) \left(\frac{dR_t}{dF_t} \frac{dF_t}{dw} + \frac{dR_t}{dF_{t-1}} \frac{dF_{t-1}}{dw} \right) \end{aligned} \quad (\text{back propagation})$$

Where $\frac{dS_T}{dA}, \frac{dA}{dR_t}, \frac{dS_T}{dB}, \frac{dB}{dR_t}$ can be obtained at each step, and

$$\frac{dR_t}{dF_t} = -u \cdot \text{sign}(F_t - F_{t-1}) \quad \frac{dR_t}{dF_{t-1}} = r_t + u \cdot \text{sign}(F_t - F_{t-1})$$

$$\frac{dF_t}{dw} = \frac{d}{dw} \{ \tanh(w^T x_t) \} = (1 - \tanh^2(w^T x_t)) \cdot (x_t + \frac{dF_{t-1}}{dw} \cdot w_{N+2})$$

$\frac{dF_t}{dw}$ is recurrent and depend on the previous $\frac{dF_{t-1}}{dw}$, which can be iteratively programmed

RRL: Implementation



- Implementation:
 - Split historical data into training and test set
 - Using training data to determine the optimal weight
 - Apply the weight on test set and observe the performance
 - Assume trade only one share per transaction



- As the learning steps increases, the “Sharp Ratio” increases

```
D:\python34\python.exe C:/Users/acer/Desktop/Presentation/rrl_stock.py
Please enter the initial weight coefficient (0.01-1.00): 0.1
Please enter the learning Speed (0.1-5.0): 2
Reinforcement Learning 1 th step, Sharp Ratio: -0.015538848940336492
Reinforcement Learning 2 th step, Sharp Ratio: 0.06929876556275599
Reinforcement Learning 3 th step, Sharp Ratio: 0.1287925043522688
Reinforcement Learning 4 th step, Sharp Ratio: 0.12907035149680088
Reinforcement Learning 5 th step, Sharp Ratio: 0.13037694126475244
Reinforcement Learning 6 th step, Sharp Ratio: 0.14789008367970302
Reinforcement Learning 7 th step, Sharp Ratio: 0.15682476764671893
Reinforcement Learning 8 th step, Sharp Ratio: 0.1506964686043493
Reinforcement Learning 9 th step, Sharp Ratio: 0.14426021998371022
Reinforcement Learning 10 th step, Sharp Ratio: 0.15453325780698657
Reinforcement Learning 11 th step, Sharp Ratio: 0.1680635929509651
Reinforcement Learning 12 th step, Sharp Ratio: 0.16540300311604716
Reinforcement Learning 13 th step, Sharp Ratio: 0.17892174648168402
Reinforcement Learning 14 th step, Sharp Ratio: 0.17956181701849072
Reinforcement Learning 15 th step, Sharp Ratio: 0.18498154040414663
```

RRL: Implementation on Stocks



- AMZN, Day-Granularity, **no loss limit control**, using 21 days price:
 - Training Data: 2013/1/2 – 2015/5/29, 606 days
 - Test Data: 2015/6/1 – 2015/9/30, 86 days
 - Training natural net return 171.08, test natural net return (nnt) 80.97
 - Initial weight 0.02, learning speed 0.3, commission 0.05, steps 1000





- Training: net return around 850, 500% nnt
- Test: net return around 220, 270% nnt
- If only hold long position net return 140, 170% nnt

RRL: Implementation on Stocks



- AAPL, Day-Granularity, **no loss limit control**, using 21 days price:
 - Training Data: 2012/7/2 – 2015/11/30, 859 days
 - Test Data: 2015/12/1 – 2016/3/31, 83 days
 - Training natural net return 33.65, test natural net return -12.24
 - Initial weight 0.1, learning speed 1.0, commission 0.05, steps 1000



- Training: net return around 197.0, 600% nnt
- Test: net return around -25.0, 200% nnt (loss)
- Biggest drawback happened at the second half of Jan, 2016, when the stock price was bumpy

RRL: Implementation on Futures





- China PE(L) Futures, Minute-Granularity
 - Point value 5, Commission fee 3 per share
 - Assume no slippage (but can be included by raising commission fee)
 - Simplified data, only use close price
 - Training Data: 2013/1/4 – 2014/8/29, 90224 tuples
 - Test Data: 2014/9/1 – 2014/11/28, 13275 tuples
- More control: decrease risk exposure by lowering trading times and forbid overnight positions, cut loss by setting absolute loss limit
 - Trading Period: 9:30 A.M. – 2:30 P.M. Overnight not allowed
 - Loss limit control: e.g. cut when sum loss greater than 200 in 3 bars
 - Lower trading frequency: Hold same position for at least 5 bars

RRL: Implementation on Futures



- Performance

- Initial weight 0.03, learning speed 0.5, steps 200, #bars 21, loss limit 200
- Training: Net profit 36409.0, 820% the original margin; Trading times 9114, around 24 transactions per day (still too high)
- Test: Net profit 3702.0, 86% the original margin; Trading times 1357, around 20 transactions per day





RRL: Implementation on Futures



- Performance

- Initial weight 0.02, learning speed 0.8, steps 300, #bars 30, loss limit 175
- Training: Net profit 28783.0, 650% the original margin; Trading times 9427, around 25 transactions per day (still too high)
- Test: Net profit 673.0, 15% the original margin; Trading times 1392, around 20 transactions per day, significant drawback





RRL: Implementation on Futures



- Performance

- Initial weight 0.1, learning speed 1.0, steps 300, #bars 14, loss limit 150
- Training: Net profit 34424.0, 650% the original margin; Trading times 9096, around 24 transactions per day (still too high)
- Test: Net profit 5331.0, 124% the original margin; Trading times 1349, around 20 transactions per day



RRL: Conclusion



- Conclusion
 - Training by RRL, the utility function (S_T) can be considerably optimized
 - The raw algorithm shows promising result on test data, generally it will perform better than the natural return
 - The strategy developed by RRL is likely to be a trend following strategy, for it performs worse when it is bumpy, and much better when there is a good trend
 - The style of the strategy on higher frequency data is like “hit and run”, very swift and sensitive. Under minute-granularity, we have to filter the trading signals to lower the trading frequency, thus avoiding slippage.
 - The strategy with such trading style may perform better on even higher frequency data, e.g. High frequency trading on Forex under tick-level.

RRL: Future Work



- Improvement
 - Improve loss control mechanism (e.g. trailing stop loss limit)
 - Use cross validation to split, train, and test.
 - Optimize the parameters other than weight. e.g. #bars
 - Revise the benchmark metrics, include in more factors. e.g. Trading times
 - Test the feasibility on tick-level forex trading